


Visualizing Memory

- Logical view of memory is a 2-dimensional matrix
 - The smallest addressable unit is the byte (8 bits)
 - We also need to be able to address halfwords (16 bits), words (32 bits) and double words (64 bits)

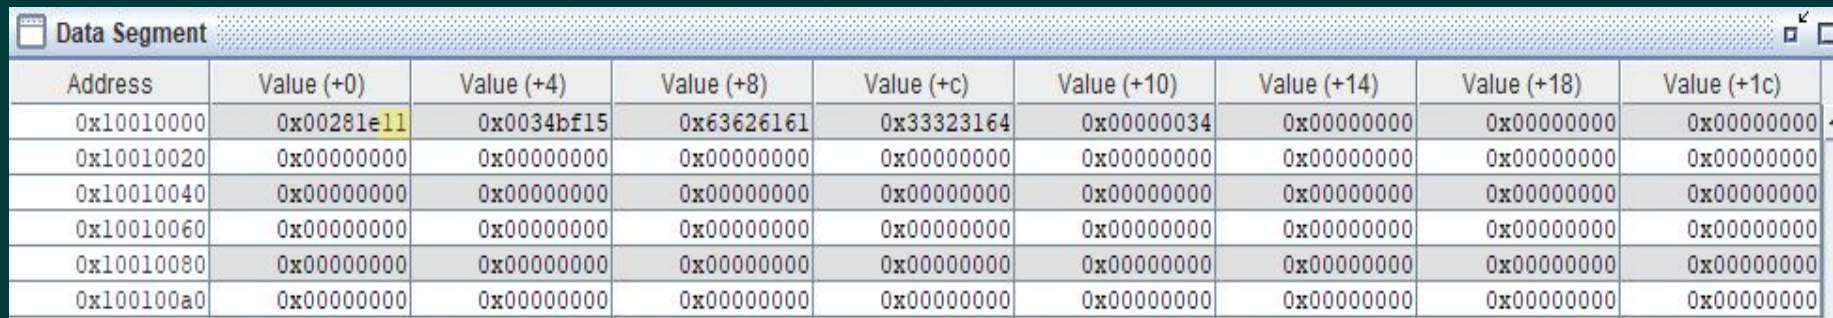


- Each cell stores 1 byte
- Each row stores 1 word
- Two rows are required for double words
- Memory addresses can be used to reference a byte, a halfword, word or double word
 - The size of data to be referenced is defined in the memory reference instruction (more on this to come).

RARS Programming View of Memory

- The programming view of memory is organized by word groups
 - 4-byte groups
 - Rows and columns

This is what memory will look like when we are programming in RARS.




The screenshot shows a window titled "Data Segment" with a table of memory addresses and their corresponding values. The table has 9 columns: Address, Value (+0), Value (+4), Value (+8), Value (+c), Value (+10), Value (+14), Value (+18), and Value (+1c). The first row shows a non-zero value at address 0x10010000, while all other rows show zero values. An arrow points from the text box above to the first row of the table.

| Address | Value (+0) | Value (+4) | Value (+8) | Value (+c) | Value (+10) | Value (+14) | Value (+18) | Value (+1c) |
|------------|------------|------------|------------|------------|-------------|-------------|-------------|-------------|
| 0x10010000 | 0x00281e11 | 0x0034bf15 | 0x63626161 | 0x33323164 | 0x00000034 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010020 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010040 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010060 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010080 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x100100a0 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |

- Each row contains 8 word cells – 4 bytes each cell
- Bytes are ordered “little endian” – low order address is rightmost byte within word
 - First data item is hex 11 shown in yellow highlighting at address 0x10010000
 - Subsequent bytes from right to left, then continue in next cell
 - Columns label offset of low order byte in each cell

Logical View vs. RARS View



← 0x1001000c
← 0x10010008
← 0x10010004
← 0x10010000

| Data Segment | | | | | | | | |
|--------------|------------|------------|------------|------------|-------------|-------------|-------------|-------------|
| Address | Value (+0) | Value (+4) | Value (+8) | Value (+c) | Value (+10) | Value (+14) | Value (+18) | Value (+1c) |
| 0x10010000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010020 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010040 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010060 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x10010080 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x100100a0 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x100100c0 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |
| 0x100100e0 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 | 0x00000000 |

←

→

0x10010000 (.data)

☒ Hexadecimal Addresses

☒ Hexadecimal Values

☐ ASCII

Assembly Language


- Hardware level programming language
 - Each instruction equates to one machine operation
- Instructions are executed sequentially
 - Each instruction is read from memory, decoded and executed (abstract view – more details in Ch. 4)
- The operations defined in a single HLL statement may result in several individual machine operations
- Assembly language has a rigid syntax because it is designed for specific hardware

Instructions


- **Instruction** is the basic unit of work in the CPU.
- Represents a single task (**operation**) to perform.
- Instructions have two parts:
 - The **operation mnemonic** (every instruction has one)
 - The **operands** to be used in performing the operation
 - Exceptions: Not all instructions need data – most do
- Types of operations
 - **Memory access** (data transfers - reading/writing data to/from memory)
 - **Arithmetic/logic** (add, subtract, and, or, etc.)
 - **Control** (decision making – branching, looping)
 - **System calls** (call on OS to perform a function)

Program Patterns in Assembly Language

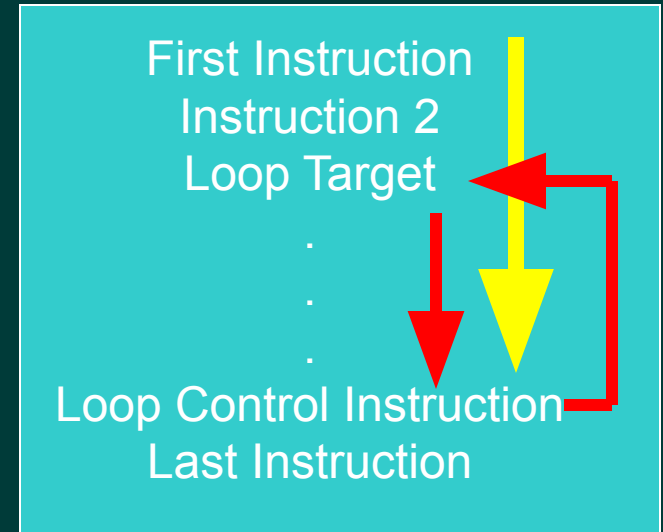
Sequential (Straight Line Code)




Branching (Skip Over Code)



Looping




Subroutine Calls



RISC-V Assembly Language

Program Structure



`main:` and `exit:` are labels. Labels define references to locations in memory. A memory address is associated with every label by the assembler.

} Define data (data segment)

} Perform operations on data (text segment)

`.data` and `.text` are *assembler directives*.

Assembler directives begin with a period and are used to convey information to the assembler and tell it how to translate an assembly language program to machine code.

Data Types

- Refers to a classification of data that specifies what kind of information a variable can hold
 - Usually determines the operations that can be performed on it
 - And how the data is stored in memory
 - It tells the computer how to interpret the value of a piece of data
- Primitive
 - Built-in, intrinsic, basic, single values
- Non-primitive
 - Data types that store a collection of values in various formats, rather than just a single value

Java Data Types

byte (p) 8-bit signed int
short (p) 16-bit signed int
int (p) 32-bit signed int
long (p) 64-bit signed
int
char (p) 16-bit Unicode

String (n-p) class

p = primitive

n-p = non-primitive
(all statically typed)

RISC-V Data Elements

byte 8-bit signed int
half 16-bit signed int
word 32-bit signed int
dword 64-bit signed int
ascii 8-bit UTF-8
Unicode

Assembly level
primitives;
not statically typed; use
is dependent on context

asciz null-term char

Defining Data in Memory

- RISC-V **assembler directives** in .data segment
 - .byte - 8-bit numeric value
 - .half - 16-bit numeric value
 - .word- 32-bit numeric value
 - .dword - 64-bit numeric value
 - .ascii - ASCII character (1 or more)
 - .asciz - null terminated ASCII character (1 or more)
 - .space - specify number of bytes to allocate
- Most data items are labeled for reference (symbol name)
 - Synonymous with user-defined variable names in HLLs
 - Symbol names must be unique
 - Can't use same label for multiple data items
- Data items should be initialized to some value
- Multiple values can be allocated with a single directive
 - Ex. arrays (more on this later)

Example Data Definition in RISC-V

| <u>Code</u> | <u>Comments</u> |
|--------------------------|------------------------------------|
| .data | # begins data segment |
| num: .byte 17 | # 8-bit number |
| array: .byte 30 40 50 60 | # four 8-bit numbers (array) |
| lrgnum: .half 600 | # 16-bit number |
| bignum: .word 3456789 | # 32 bit number |
| char: .ascii "a" | # single character |
| strng: .asciz "abcd1234" | # null terminated string (9 bytes) |
| extra: .space 10 | # memory space for 10 bytes |

Notes:

1. Each data value has a symbol name. Symbol names are user defined (just like variable names in HLLs).
2. Multiple values can be delimited by a blank or by a comma (arrays).
3. Characters and strings must be enclosed in double quotation marks.
4. The extra definition allocates 10 bytes of memory. Anything can be stored in this space (10 bytes or character data).
5. The numeric values specified are by default base 10. You can also specify hexadecimal values by using the prefix 0x and the hex number.

RARS Memory Allocation

```
.data
num:   .byte   17
array: .byte   30 40 50 60
lrgnum: .half 600
bignum: .word   3456789
char:   .ascii  "a"
strng:  .asciz  "abcd1234"
extra:  .space  10
```

| | | | | |
|-----------------|-----|--------|-----|------------|
| space (4 bytes) | | | | 0x1001001c |
| space (4 bytes) | | | | 0x10010018 |
| space (2 bytes) | | "null" | "4" | 0x10010014 |
| "3" | "2" | "1" | "d" | 0x10010010 |
| "c" | "b" | "a" | "a" | 0x1001000c |
| 3456789 | | | | 0x10010008 |
| 600 | | ⊗ | 60 | 0x10010004 |
| 50 | 40 | 30 | 17 | 0x10010000 |
| 3 | 2 | 1 | 0 | |

| Address | Value (+0) | Value (+4) | Value (+8) | Value (+c) | Value (+10) | Value (+14) | Value (+18) | Value (+1c) |
|------------|------------|------------|------------|------------|-------------|-------------|-------------|-------------|
| 0x10010000 | 0x32281e11 | 0x0258003c | 0x0034bf15 | 0x63626161 | 0x33323164 | 0x00000034 | 0x00000000 | 0x00000000 |

| Label | Address ▲ |
|---------------------------|------------|
| Data SegmentExampleFro... | |
| num | 0x10010000 |
| array | 0x10010001 |
| lrgnum | 0x10010006 |
| bignum | 0x10010008 |
| char | 0x1001000c |
| strng | 0x1001000d |
| extra | 0x10010016 |

Empty memory location

- Memory hole
- Could not be filled due to data values defined & sequencing

Alignment Restriction

- Defines where numeric values are placed when defined in memory
- Offset restriction based on size of data
- RISC-V default:
 - Halfword values must begin at memory addresses that are a multiple of 2
 - Word values must begin at memory addresses that are a multiple of 4
 - Double-word values must begin at memory addresses that are a multiple of 4

Memory Offset

- The memory offset is a value that is added to the base address forming the actual memory address of the memory location to be referenced
- The offset value is determined by the order and types of data defined in the data segment
- Example: our previous data definition

```
.data
num:    .byte    17  ← offset 0 (always for 1st data value)
array:  .byte    30 40 50 60 ← offset 1 (30), 2 (40), 3 (50), 4 (60)
lrgnum: .half    600 ← offset 6 (extending through 7)
bignum: .word    3456789 ← offset 8 (extending through 11)
char:   .ascii   "a" ← offset 12
strng:  .asciz   "abcd1234" ← offset 13 (extending through 21)
extra:  .space   10 ← offset 22 (extending through 31)
```

Ponder This ...

.data

| | | | |
|--------|--------|------------|---------------------------|
| one: | .byte | 97 98 99 | # 3-element array |
| two: | .word | 0x00636261 | # int defined in hex |
| three: | .asciz | "abc" | # 3-char null term string |
| four: | .word | 6513249 | # int defined in decimal |

In memory:

