

Computer Organization and Design

Introduction to Computer Performance
(Sections 1.6 - 1.10 from Chapter 1
Plus Ancillary Information)

Understanding Performance

- Algorithm
 - Determines number of operations executed
- Programming language, compiler, architecture
 - Determines number of machine instructions executed per operation
- Processor and memory system
 - Determines how fast instructions are executed
- I/O system (including the operating system)
 - Determines how fast I/O operations are executed

Performance Goals

- Designing high performance computers is an important goal of a computer architect
 - Cost vs. performance
- Assessing the performance of computer hardware is an important part of computer design
 - Affects the demand and market value of the computer
- How to quantify computer performance
 - What application to use to measure performance?
 - What component of computer to measure? (CPU, I/O, memory)
 - How do other parameters affect performance? (OS, compiler)
 - How do you define performance – faster CPU/memory or quantity of work completed per unit time?

Executing a Program

- Load the binary executable (or a minimum portion) into memory
 - Some I/O occurs to access the stored program on disk and initiates a read of the program file
 - Operating system allocates memory space for the program
- Operating system initializes CPU for program execution
- CPU begins reading program instructions from memory and executes
- Periodic pauses in execution
 - Disk I/O – either data or more of the program is needed in memory
 - Operating System – temporarily suspends program to deal with other software or hardware operations and then resumes
- Program finishes execution
 - Operating system cleans up
 - Prepares for next program
- Time is associated with all of these operations
 - How do you narrow down what to measure?

Basic Performance Measures

- Execution time (response time)
 - Time between start and completion of task
 - Task can be any unit of work – program, job, procedure
 - Focus is on time
- Throughput
 - Total amount of work completed in a given time
 - Jobs/day, processes/hour, instructions/second
 - Focus is on work
- Importance of one over the other depends on perspective – ex. user vs CIO

Abstract Example

- Consider the impact of the following changes to a computer system with regard to execution time and throughput
 - 1. Replace the processor with a faster version
 - 2. Add additional processors to a computer that uses multiple processors for separate tasks
- Option 1 would improve both execution time and throughput
 - Decreasing execution time almost always improves throughput
- Option 2 would only improve throughput
 - Individual tasks execute at the same speed but more tasks can run concurrently

Performance Measures

- The primary measure of performance is execution time
 - Exception is I/O where throughput is important
- To maximize performance of an application, we need to minimize its execution time
- The relationship between performance and execution time is represented by the expression

$$Performance_x = \frac{1}{Execution Time_x}$$

Performance Comparison

- When comparing two systems, we can say if the performance of computer X is better than computer Y, then

$$Performance_X > Performance_Y$$

$$Execution\ Time_Y > Execution\ Time_X$$

$$\frac{1}{Execution\ Time_X} > \frac{1}{Execution\ Time_Y}$$

Performance Comparison (continued)

- The quantitative difference in performance between two computers is represented

$$\frac{Performance_x}{Performance_y} = \frac{Execution\ time_y}{Execution\ time_x} = n$$

- X is n times faster than Y
- If X is n times faster than Y, then the execution time on Y is n times longer than it is on X
- The value n is the ratio of the measured performance between two computer systems

Performance Comparison (continued)

- Example:
 - If computer A runs a program in 10 microseconds and computer B runs the same program in 15 microseconds, how much faster is A than B?

$$\frac{\textit{Execution time}_B}{\textit{Execution time}_A} = \frac{15}{10} = 1.5$$

System Performance

- Total system performance depends on the performance of the individual components:
 - CPU
 - Memory
 - I/O
- CPU – execution engine
 - Performance defined by computational speed
- Memory – short term storage
 - Performance defined by speed of access
- I/O - long term storage
 - Performance defined by data throughput

Measuring Performance

- Time is the measure of computer performance
- Two kinds of time
 - Wall-clock time (**response time**, elapsed time)
 - used to measure the total time to complete a task or job
 - includes I/O time, memory accesses, and OS overhead
 - CPU execution time (**CPU time**)
 - used to measure the time the processor spends working on a single task or job
 - not including I/O time or time spent on other tasks
- CPU time can be further divided
 - **User CPU time** – time spent working on a program
 - **System CPU time** – time spent by OS doing things for the program



Measuring Performance (continued)


- CPU performance measurements
 - use CPU time rather than response time because it is more accurate
- User CPU time is preferred
 - because of its independence of the operating system and other factors
- The time required for a computer to perform a task (run a program) is dependent on the speed of the hardware it runs on
 - Execution time is measured in seconds per program

Finite State Machine

- A mathematical model of computation used to design both computer programs and sequential logic circuits
- A computer is an implementation of a finite state machine
 - At any given point in time, the state of the machine is defined by the individual values present on the inputs and outputs of its components
- These values are combinations of 0s and 1s
- The operation of the machine consists of transitions from one state to the next
- The change in state is initiated by a triggering event
 - The clock

The Clock

- Clock signals



- Clock frequency

- Cycles per second (CPS) or Hertz (Hz) [kHz, MHz, GHz]

- Clock period

- Seconds or fraction of (millisecond, microsecond, nanosecond)
- Calculated as inverse of frequency ($1/\text{Hz}$)

- Rise time/Fall time

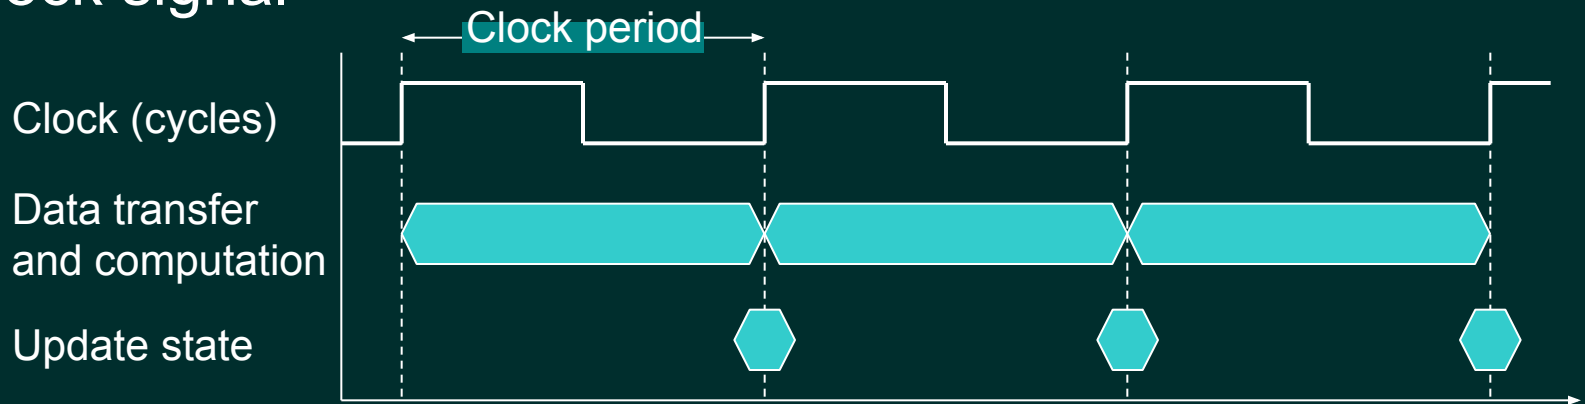
- Time to go from 0 to 1 or from 1 to 0

Actual Clock Waveform



Hardware and Timing

- A circuit called a *clock* generates timing signals that control all the other circuits in the computer
- Clock signal



- A computer's clock speed is described using two terms
 - Frequency - cycles per second or Hertz (Hz)
 - **clock rate**
 - Period - time for one cycle – inverse of frequency - seconds
 - **clock cycle time**

Calculating CPU Time

- CPU time is calculated using the following two formulas

$$CPU_{Time} = CPU \text{ Clock Cycles} \times \text{Clock Cycle Time}$$

$$CPU_{Time} = \frac{CPU \text{ Clock Cycles}}{Clock \text{ Rate}}$$

- CPU performance is improved
 - By reducing the length of a clock cycle or
 - By reducing the number of clock cycles required for a program

Example

- CPU A
 - Runs a program in 10 seconds
 - Has a 4 GHz clock
- CPU B
 - Being designed
 - Want to run the same program in 6 seconds
 - The clock rate can be made to run faster
 - BUT it will affect the CPU design such that the new computer will need 1.2 times as many clock cycles as CPU A.
- What clock rate do we need to achieve on the new CPU?

Example (continued)

- First, calculate the number of clock cycles required for the program on CPU A

$$CPU\ time_A = \frac{CPU\ clock\ cycles_A}{Clock\ rate_A}$$

$$\begin{aligned} CPU\ clock\ cycles_A &= CPU\ time_A \times Clock\ rate_A \\ &= 10\ seconds \times (4 \times 10^9)\ cycles\ per\ second \\ &= 40 \times 10^9\ cycles \end{aligned}$$

Example (continued)

- The CPU time for CPU B can now be calculated as follows

$$CPU\ time_B = \frac{1.2 \times CPU\ clock\ cycles_A}{Clock\ rate_B}$$

$$6\ seconds = \frac{1.2 \times 40 \times 10^9\ cycles}{Clock\ rate_B}$$

$$Clock\ rate_B = \frac{1.2 \times 40 \times 10^9\ cycles}{6\ seconds}$$

$$Clock\ rate_B = 8\ GHz$$

- CPU B needs a clock rate twice that of CPU A
 - But can we actually achieve that clock rate?
 - Are there other issues that need to be considered?

Program Execution Time

- The CPU time for a program directly depends on the number of instructions in that program

$$\text{CPU clock cycles} = \text{\# instructions for a program} \times \text{Average clock cycles per instruction}$$

- The term *clock cycles per instruction* is often abbreviated CPI
- The CPI of a program depends on the instruction set of the computer and the compiler
- In other words, the CPI depends on the architecture of the system

Example

- Suppose we have two implementations of the same instruction set architecture.
 - Example: Intel x86 vs. AMD x86
 - CPU A has a clock cycle time of 10 ns and a CPI of 2.0 for a given program.
 - CPU B has a clock cycle time of 20 ns and a CPI of 1.2 for the same program.
- Which machine is faster for this program?
- Assume the program requires I instructions to be executed.
 - I is the variable representing the unknown number of instructions

Example (continued)

- First calculate the clock cycles for each CPU
 - CPU clock cycles A = $I \times 2.0$
 - CPU clock cycles B = $I \times 1.2$
- Next calculate CPU time
 - CPU time A = $2.0I \times 10 \text{ ns} = 20I \text{ ns}$
 - CPU time B = $1.2I \times 20 \text{ ns} = 24I \text{ ns}$
- CPU A is $24 / 20 = 1.2$ times faster than CPU B

Program CPU Time

- The CPU time for a program, which is our main measure of performance, can be written as

$$\text{CPU time} = \frac{\text{Instruction count}}{\text{CPI}} \times \text{Clock Cycle time}$$

$$\text{CPU time} = \frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}$$

- The performance of the CPU is directly dependent on
 - the clock speed
 - the number of cycles per instruction
 - the number of instructions per program - instruction count (IC)

CPI (Cycles Per Instruction)

- An instruction is a unit of work in the CPU
- Different types of instructions require more or less time to execute in terms of number of cycles
 - Ex. An add instruction may require fewer clock cycles to execute than a memory access instruction
- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- Weighted average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \underbrace{\frac{\text{Instruction Count}_i}{\text{Instruction Count}}}_{\text{Relative frequency}} \right)$$

Relative frequency

Calculating CPI

- A RISC-V microprocessor has been designed to achieve a target clock frequency. The various instruction classes require varying clock cycles as shown in the following table:

Intstruction Class	Instruction Latency
Loads	5 cycles
Stores	4 cycles
Branches	3 cycles
Jumps	2 cycles
ALU/Logic	4 cycles

- What is the average CPI for the microprocessor?
 - Simply average the latencies for all instruction classes
 - $CPI = (5 + 4 + 3 + 2 + 4) / 5 = 3.6$ cycles

Calculating CPI

- A program is executed on the previously defined RISC-V microprocessor. The relative frequency of each class of instruction, which is program dependent, is added to the table:

Intstruction Class	Instruction Latency	Relative Frequency
Loads	5 cycles	30%
Stores	4 cycles	15%
Branches	3 cycles	10%
Jumps	2 cycles	5%
ALU/Logic	4 cycles	40%

- What is the average CPI for the program?
 - Factor the instruction frequency with instruction latency
 - $$\text{CPI} = (5 * 0.3) + (4 * 0.15) + (3 * 0.1) + (2 * 0.05) + (4 * 0.4)$$
$$= 4.1 \text{ cycles}$$

Components of Performance

Performance Component	Unit of Measure
CPU execution time for a program	Seconds for the program to run
Instruction count	Instructions executed for the program
Clock cycles per instruction (CPI)	Average number of clock cycles per instruction
Clock cycle time	Seconds per clock cycle

Contributors to Performance

Hardware or Software Component	Affects What?
Algorithm	Instruction count and possibly the CPI
Programming language	Instruction count and CPI
Compiler	Instruction count and CPI
Instruction set architecture	Instruction count, clock rate and CPI

Performance Tradeoffs

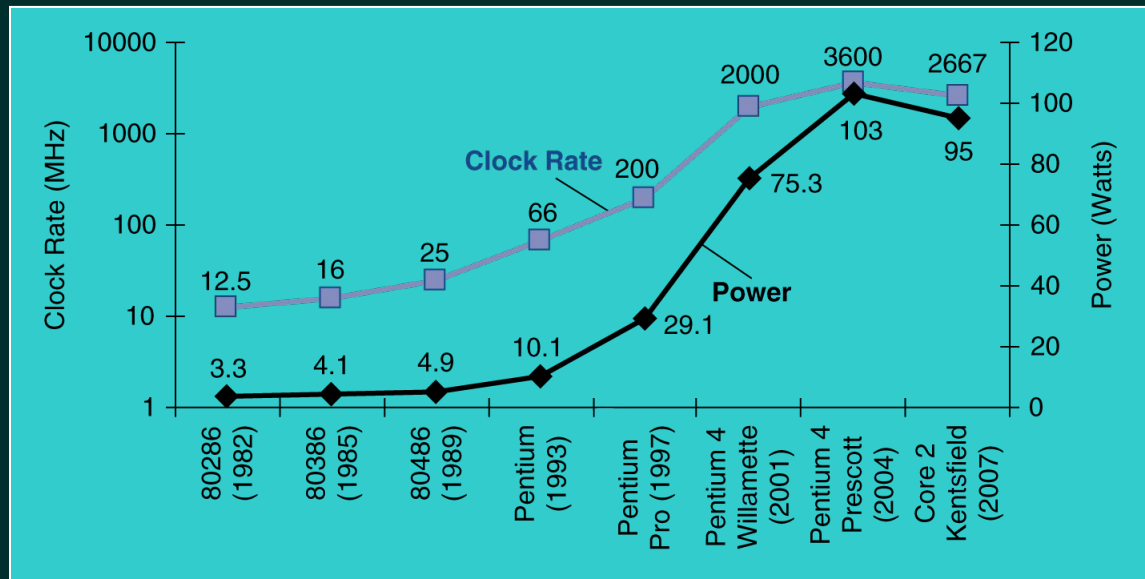
- Performance is equally dependent on each of the three terms: IC, CPI, and clock cycle time
- When comparing machines, you must consider all three components
- There are trade-offs between the factors that affect performance
 - optimizing one of them can lead to worsening the performance of another
 - degrade overall performance

Benchmarks

- Programs that are used to measure performance
- Execute a specific set of programs that represents the typical real world workload
- Most widely used benchmark is SPEC (Standard Performance Evaluation Corporation)
 - CPU
 - Graphics/Workstations
 - High-Performance Computing, MPI/OMP
 - Java Client/Server
 - Mail Servers
 - Network File System
 - Power
 - SIP (Session Initiation Protocol)
 - Virtualization
 - Web Servers

Power Trends

- Dominant technology for integrated circuits today is CMOS (Complimentary Metal Oxide Semiconductor)
 - Primary power consumed during switching (dynamic power)
- Dynamic power is measured in watts based on the following equation:
 - $\text{Power}_{\text{Dynamic}} = 1/2 \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Switching frequency}$



Clock rate and power for Intel x86 microprocessors over 8 generations

Power Problems

- Want to build faster processors?
 - Natural tendency is to increase frequency which makes power increase.
 - So, now reduce voltage to help reduce power but find that more power is consumed by leaky transistors (even the ones that are turned off and not switching).
 - So, increase insulation properties of transistors but it only helps a little and there's not much room to add a lot of insulation around the transistors and wires and still get everything to fit on the chip.
- Ultimately, you get to the point where you can't go forward without employing more expensive cooling methods that are not practical for consumer systems.
- This limitation is referred to as the **power wall**
 - Can't improve performance with current designs by increasing frequency, lowering voltage and reducing capacitive load
 - Limit for cooling commodity general purpose processors

The Shift to Multiprocessors

- Small increases in performance in uniprocessors (single core) will continue to be made but not at the same rate as in the past.
 - There is still a market for uniprocessors (embedded systems/IoT)
- New general purpose CPU designs are based on multicore processors
 - More than one CPU (core) on a single chip
 - Newer designs are hybrid core (big.LITTLE or other)
- New Moore's Law?
 - What will define and guide future development of microprocessors when doubling transistor counts no longer can be done?
 - Domain specific architectures, more hybrid designs
- Not feasible to continue increasing core counts when most applications cannot utilize more processors
 - There is a need to program differently to take advantage of multiple hardware

Parallelism

- Parallelism is a concept that describes the ability to run more than one thing at a time.
 - Multiprocessor mainframes have existed since the 1970s
 - Each processor (a uniprocessor) could be assigned a program or a part of a program to run under operating system control
 - Goal was to increase throughput
- Instruction level parallelism (ILP) refers to the abstract parallelism among instructions in a program
 - If two instructions are not dependent on each other, then they could be executed in parallel providing adequate hardware resources
 - Pipelining (Chapter 4) is a form of ILP
 - Hardware executes multiple instructions at once
 - Details are hidden from programmer
 - Don't have to change or recompile programs for ILP

Multiprocessor Performance

- Gains in performance for multicore processors requires programming for performance
 - Programmers have to design their programs to take advantage of parallelism
 - Explicitly Parallel Programming (EPP)
 - Can't rely on hardware alone to increase performance
 - Load balancing becomes programmer responsibility
 - Inter-processor communication and synchronization must be optimized
 - New programming paradigm
 - Programmers have to learn new ways of writing algorithms to achieve higher parallelism

Amdahl's Law

- Used to calculate improvement in performance based on isolating the feature improved.
- A rule stating that the performance enhancement possible with a given improvement is limited by the amount that the improved feature is used.
- Also known as the law of diminishing returns.
- Is used to calculate improvements in performance for both software and hardware

Amdahl's Law (Example 1)

- Suppose we are considering an enhancement that runs 10 times faster than the original system but is used only 40% of the time. What is the overall speedup gained by incorporating the enhancement?

- Amdahl's Law formula:

$$S_o = \frac{1}{(1 - F) + \frac{F}{S_E}}$$

- S_E = speedup of enhancement
- F = fraction enhancement is used
- S_o = overall speedup

$$S_o = \frac{1}{(1 - 0.4) + \frac{0.4}{10}} = 1.5625$$

Amdahl's Law (Example 2)

- A program takes 100 seconds to run; 80 seconds are consumed by multiply instructions. How much do I have to improve the multiplication speed so the program will run 5 times faster?

Execution time after improvement =

$$\frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}}$$

+ Execution time unaffected

Amdahl's Law Example (cont)

Execution time after improvement

$$= \frac{80 \text{ seconds}}{n} + (100 - 80 \text{ seconds})$$

$$20 \text{ seconds} = \frac{80 \text{ seconds}}{n} + 20 \text{ seconds}$$

$$0 = \frac{80 \text{ seconds}}{n}$$

Can't be done!

Another Example

- If 30% of the execution time of a program can be enhanced, the speedup p will be 0.3.
- If the improvement makes the affected part twice as fast, s will be 2. Amdahl's law states that the overall speedup of applying the improvement will be:

$$S_{latency} = \frac{1}{1 - p + \frac{p}{s}} = \frac{1}{1 - 0.3 + \frac{0.3}{2}} = 1.18$$

Another Example

- Assume that we are given a serial task which is split into four consecutive parts
- The percentages of execution time are
 - $p1 = 0.11$, $p2 = 0.18$, $p3 = 0.23$, $p4 = 0.48$
- The speedups for each part is as follows:
 - $s1 = 1$, $s2 = 5$, $s3 = 20$, $s4 = 1.6$
- By using Amdahl's law, the overall speedup is

$$S_{latency} = \frac{1}{\frac{p1}{s1} + \frac{p2}{s2} + \frac{p3}{s3} + \frac{p4}{s4}} = \frac{1}{\frac{0.11}{1} + \frac{0.18}{5} + \frac{0.23}{20} + \frac{0.48}{1.6}}$$

$$S_{latency} = 2.19$$

Notice how the 5 times and 20 times speedup on the 2nd and 3rd parts respectively don't have much effect on the overall speedup when the 4th part (48% of the execution time) is accelerated by only 1.6 times.

Terminology / Vocabulary

- Execution time (response time)
- Throughput
- User CPU time vs. System CPU time
- Clock
- Frequency (clock rate)
- Period (cycle time)
- CPI (cycles per instruction)
- Benchmarks
- Power
- ILP (Instruction Level Parallelism)
- Amdahl's Law