

✓ Requirements and questions

Exploratory Analysis(10 points)

- Are there any correlation between features? why or why not?
- Are there any outliers in the dataset? why or why not?
- Are there any normalization needed? why or why not?
- Are there any missing values? Why or why not?

Classification (25 points)

- Experiment the following algorithms/models: decision tree, random forest, adaboost, KNN, SVM, MLP, and Naive Baye
- For each model, train-test-split with 80% for train and 20% for test.
- For each algorithm, output the following performance measures: accuracy, precision, and recall

Visualization (5 points)

Pick at least one visualization for model performance comparison.

Reflection (15 points)

- What is the best K in KNN? Only consider k in the range of 1-15. Use odd numbers only.
- What is the most important features found by decision tree and random forest?
- Which algorithm has the highest accuracy? Is there model overfitting for this algorithm? Why or why not? Perform 10-fold cross-validation with this algorithm and report the accuracy, precision and recall.

```
#pasted from URL
from ucimlrepo import fetch_ucirepo

# fetch dataset
iris = fetch_ucirepo(id=53)

# data (as pandas dataframes)
X = iris.data.features
y = iris.data.targets

# metadata
print(iris.metadata)

# variable information
print(iris.variables)
```

```
{'uci_id': 53, 'name': 'Iris', 'repository_url': 'https://archive.ics.uci.edu/dataset/53/iris'}

name      role      type      demographic \
0  sepal length  Feature  Continuous      None
1  sepal width   Feature  Continuous      None
2  petal length  Feature  Continuous      None
3  petal width   Feature  Continuous      None
4      class     Target  Categorical      None

description units missing_values
0              None      cm      no
1              None      cm      no
2              None      cm      no
3              None      cm      no
4  class of iris plant: Iris Setosa, Iris Versico...  None      no
```

```
correlation_matrix = X.corr(method='pearson')
print(correlation_matrix)
```

```
#yes there is a strong correlation between the pedal length
```

	sepal length	sepal width	petal length	petal width
sepal length	1.000000	-0.109369	0.871754	0.817954
sepal width	-0.109369	1.000000	-0.420516	-0.356544
petal length	0.871754	-0.420516	1.000000	0.962757
petal width	0.817954	-0.356544	0.962757	1.000000

```
import numpy as np

z_scores = np.abs((X - X.mean()) / X.std())
outliers = (z_scores > 3).any(axis=1)
print(X[outliers])
print("num of outliers:", outliers.sum())

#there is one outlier present this was found ing z_score to detmine the
```

```
    sepal length  sepal width  petal length  petal width
15             5.7         4.4           1.5           0.4
num of outliers: 1
```

```
# check for missing vals
dups = X.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))

# yes we have 3 duplicate rows found in hte X set the .duplicated() func
```

```
Number of duplicate rows = 3
```

```
print(X.isnull().sum())
# no there are no null vlaues we know this because of the function call w
```

```
sepal length    0
sepal width     0
petal length    0
petal width     0
dtype: int64
```

✓ part 2: Classification

```
#global imports and data split
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, StratifiedKFold, c
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, precision_score, recall_score

# im just using the same train test split for each model

# y -> 1D numeric
y_1d = np.asarray(y).squeeze()
le = LabelEncoder()
y_enc = le.fit_transform(y_1d)

# split
X_train, X_test, y_train, y_test = train_test_split(X, y_enc, test_size=
```

```
#decision tree classifier

# train
clf = tree.DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

# eval
y_pred = clf.predict(X_test)
print(f"accuracy: {accuracy_score(y_test, y_pred):.3f}")
print(f"precision (macro): {precision_score(y_test, y_pred, average='mac
print(f"recall (macro): {recall_score(y_test, y_pred, average='macro'):.

# visualization of each level of decision tree
plt.figure(figsize=(12,8))
tree.plot_tree(clf, feature_names=X.columns, class_names=sorted(y_1d.uniq
plt.show()

importances = pd.Series(clf.feature_importances_, index=X.columns).sort_
print("DTfeature importances:\n", importances)
```

```
accuracy: 0.933
precision (macro): 0.933
```

```
graph TD
    Node0["petal length <= 2.45  
gini = 0.667  
samples = 120  
value = [40, 40, 40]  
class = Iris-setosa"] --> Node1["petal width <= 1.65  
gini = 0.5  
samples = 80  
value = [0, 40, 40]  
class = Iris-versicolor"]
    Node0 --> Node2["petal length <= 4.95  
gini = 0.133  
samples = 42  
value = [0, 39, 3]  
class = Iris-versicolor"]
    Node1 --> Node3["petal length <= 4.85  
gini = 0.051  
samples = 38  
value = [0, 1, 37]  
class = Iris-virginica"]
    Node1 --> Node2
    Node2 --> Node4["sepal length <= 6.15  
gini = 0.375  
samples = 4  
value = [0, 1, 3]  
class = Iris-virginica"]
    Node2 --> Node5["gini = 0.0  
samples = 38  
value = [0, 38, 0]  
class = Iris-versicolor"]
    Node4 --> Node6["sepal width <= 2.45  
gini = 0.5  
samples = 2  
value = [0, 1, 1]  
class = Iris-versicolor"]
    Node4 --> Node7["gini = 0.0  
samples = 2  
value = [0, 0, 2]  
class = Iris-virginica"]
    Node6 --> Node8["gini = 0.0  
samples = 1  
value = [0, 0, 1]  
class = Iris-virginica"]
    Node6 --> Node9["gini = 0.0  
samples = 1  
value = [0, 1, 0]  
class = Iris-versicolor"]
    Node3 --> Node10["sepal width <= 3.0  
gini = 0.444  
samples = 3  
value = [0, 1, 2]  
class = Iris-virginica"]
    Node3 --> Node11["gini = 0.0  
samples = 35  
value = [0, 0, 35]  
class = Iris-virginica"]
    Node10 --> Node12["gini = 0.0  
samples = 2  
value = [0, 0, 2]  
class = Iris-virginica"]
    Node10 --> Node13["gini = 0.0  
samples = 1  
value = [0, 1, 0]  
class = Iris-versicolor"]
```

The decision tree starts at the root node, which splits on 'petal length <= 2.45'. This results in two branches: one leading to a node where 'petal width <= 1.65' is used for splitting, and another leading directly to a leaf node classifying 'Iris-versicolor'. The 'petal width' branch further splits on 'petal length <= 4.85', leading to either 'Iris-virginica' or back to 'Iris-versicolor'. From the 'petal length <= 4.95' node, it splits on 'sepal length <= 6.15', leading to either 'Iris-virginica' or a leaf 'Iris-versicolor'. The 'sepal length <= 6.15' branch continues to split on 'sepal width <= 2.45', eventually leading to leaf nodes for both 'Iris-virginica' and 'Iris-versicolor'. Another path from the 'petal width <= 1.65' branch leads to 'sepal width <= 3.0', which also splits into leaf nodes for both classes.

```
DTfeature importances:
  petal length      0.558568
petal width        0.406015
sepal width        0.029167
sepal length       0.006250
dtype: float64
```

```
# random forrest classsifier
```

```

rf_clf = RandomForestClassifier(n_estimators=200, random_state=42)
rf_clf.fit(X_train, y_train)
y_pred = rf_clf.predict(X_test)

# metrics
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average="macro", zero_division=0)
rec = recall_score(y_test, y_pred, average="macro", zero_division=0)
print(f"Accuracy: {acc:.3f} | Precision (macro): {prec:.3f} | Recall (macro): {rec:.3f}")

print(classification_report(y_test, y_pred))

# feature importances with names
importances = pd.Series(rf_clf.feature_importances_, index=X.columns).sort_values(ascending=False)
print("feat importances:\n", importances)

# 10-fold CV with accuracy/precision/recall for hte later question
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
scores = cross_validate(
    rf_clf, X, y, cv=cv,
    scoring={"acc": "accuracy", "prec": "precision_macro", "rec": "recall_macro"}
)

```

Accuracy: 0.900 | Precision (macro): 0.902 | Recall (macro): 0.900

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	10
1	0.82	0.90	0.86	10
2	0.89	0.80	0.84	10

accuracy			0.90	30
macro avg	0.90	0.90	0.90	30
weighted avg	0.90	0.90	0.90	30

feat importances:

petal length	0.453793
petal width	0.412449
sepal length	0.115873
sepal width	0.017885

dtype: float64

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier

# base learner
stump = DecisionTreeClassifier(max_depth=1, random_state=42)
adaboost = AdaBoostClassifier(estimator=stump, n_estimators=50, learning_

# evaluate / fit
adaboost.fit(X_train, y_train)
y_pred = adaboost.predict(X_test)

print(f"accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))

```

```

/usr/local/lib/python3.12/dist-packages/sklearn/ensemble/_weight_boosting
warnings.warn(
accuracy: 0.9333
[[10  0  0]
 [ 0  9  1]
 [ 0  1  9]]

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	0.90	0.90	0.90	10
2	0.90	0.90	0.90	10
accuracy			0.93	30
macro avg	0.93	0.93	0.93	30
weighted avg	0.93	0.93	0.93	30

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import StratifiedKFold, cross_validate

cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

rows = []
for k in range(1, 16, 2):
    knn = Pipeline([
        ("scaler", StandardScaler()),

```

```

        ("clf", KNeighborsClassifier(n_neighbors=k))])

# fit on train | evaluate on test
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average="macro", zero_division=0)
rec = recall_score(y_test, y_pred, average="macro", zero_division=0)
print(f"k={k} | Test Acc: {acc:.3f} | Prec(m): {prec:.3f} | Rec(m): {rec:.3f}")

# 10-fold CV on data
scores = cross_validate(
    knn, X, y_enc, cv=cv,
    scoring={"acc": "accuracy", "prec": "precision_macro", "rec": "recall_macro"})

cv_acc = scores["test_acc"].mean()
cv_prec = scores["test_prec"].mean()
cv_rec = scores["test_rec"].mean()

rows.append([k, acc, prec, rec, cv_acc, cv_prec, cv_rec])

# summary table
results_df = pd.DataFrame(rows, columns=[
    "k", "test_acc", "test_prec", "test_rec", "cv_acc", "cv_prec", "cv_rec"])

print("\nsummary (sorted by CV accuracy..):\n", results_df)

# pick best k by CV accuracy
best_row = results_df.sort_values(
    ["cv_acc", "test_acc", "k"], ascending=[False, False, True]).iloc[0]

best_k = int(best_row["k"])
print(f"\nbest k by CV accuracy: {best_k} "
      f"(CV acc={best_row['cv_acc']:.3f}, test Acc={best_row['test_acc']:.3f})")

```

```

k=1 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967
k=3 | Test Acc: 0.933 | Prec(m): 0.944 | Rec(m): 0.933
k=5 | Test Acc: 0.933 | Prec(m): 0.944 | Rec(m): 0.933
k=7 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967
k=9 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967
k=11 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967
k=13 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967
k=15 | Test Acc: 0.967 | Prec(m): 0.970 | Rec(m): 0.967

```


Summary (sorted by CV accuracy):

	k	test_acc	test_prec	test_rec	cv_acc	cv_prec	cv_rec
6	13	0.966667	0.969697	0.966667	0.973333	0.977778	0.973333
7	15	0.966667	0.969697	0.966667	0.960000	0.966032	0.960000
2	5	0.933333	0.944444	0.933333	0.960000	0.969841	0.960000
3	7	0.966667	0.969697	0.966667	0.960000	0.968254	0.960000
5	11	0.966667	0.969697	0.966667	0.960000	0.969841	0.960000
4	9	0.966667	0.969697	0.966667	0.953333	0.964286	0.953333
1	3	0.933333	0.944444	0.933333	0.946667	0.957619	0.946667
0	1	0.966667	0.969697	0.966667	0.940000	0.946587	0.940000

Best k by CV accuracy: 13 (CV Acc=0.973, Test Acc=0.967)

```

from sklearn.pipeline import Pipeline
from sklearn.svm import SVC

# SVM with scaling
svm_clf = Pipeline([("scaler", StandardScaler()),
                    ("svc", SVC(kernel="rbf", C=1.0, gamma="scale", random_state=42))])

# fit / predict
svm_clf.fit(X_train, y_train)
y_pred = svm_clf.predict(X_test)

# evaluation
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, average="macro", zero_division=0)
rec = recall_score(y_test, y_pred, average="macro", zero_division=0)

print(f"accuracy: {acc:.3f} | Precision (macro): {prec:.3f} | Recall (macro): {rec:.3f}")
print("confusion matrix:\n", confusion_matrix(y_test, y_pred))
print("classification report:\n", classification_report(y_test, y_pred, target_names=iris.target_names))

```

Accuracy: 0.967 | Precision (macro): 0.970 | Recall (macro): 0.967

Confusion matrix:

```

[[10  0  0]
 [ 0  9  1]
 [ 0  0 10]]

```

Classification report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	0.90	0.95	10
Iris-virginica	0.91	1.00	0.95	10
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

#MLP

```

from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix

```

scale to MLP

```

mlp = Pipeline([

```

```

("scaler", StandardScaler()),
("mlp", MLPClassifier(
    hidden_layer_sizes=(64, 32),    # two layers
    activation="relu",
    solver="adam",
    alpha=1e-3,
    learning_rate="adaptive",
    max_iter=500,
    early_stopping=True,
    n_iter_no_change=10,
    random_state=42))
])

# fit / evaluate
mlp.fit(X_train, y_train)
y_pred = mlp.predict(X_test)

print("accuracy:", accuracy_score(y_test, y_pred))
print("confusion matrix:\n", confusion_matrix(y_test, y_pred))
print("report:\n", classification_report(y_test, y_pred))

```

accuracy: 0.6666666666666666

confusion matrix:

```

[[10  0  0]
 [ 0 10  0]
 [ 0 10  0]]

```

report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	0.50	1.00	0.67	10
2	0.00	0.00	0.00	10
accuracy			0.67	30
macro avg	0.50	0.67	0.56	30
weighted avg	0.50	0.67	0.56	30

```

/usr/local/lib/python3.12/dist-packages/sklearn/preprocessing/_label.py:1
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.p
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.p
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.p
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

```
#Naive Baye
from sklearn.naive_bayes import GaussianNB

gnb = GaussianNB(var_smoothing=1e-9)
gnb.fit(X_train, y_train)

y_pred = gnb.predict(X_test)
print("accuracy:", accuracy_score(y_test, y_pred))
print("confusion matrix:\n", confusion_matrix(y_test, y_pred))
print("report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.9666666666666667

Confusion matrix:

```
[[10  0  0]
 [ 0  9  1]
 [ 0  0 10]]
```

Report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	0.90	0.95	10
Iris-virginica	0.91	1.00	0.95	10
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

Start coding or [generate](#) with AI.

Reflection questions:

- What is the best K in KNN? Only consider k in the range of 1-15. Use odd numbers only. The best vlaue for K is 3, 5 and 9 performed the best with perfect accuracy. Other vlaues of K are close but not as high as those three.
- What is the most important features found by decision tree and random forest? from randmforrest the most important feature is: `petal` and for hte decision tree the most important feature is `petal length` with the width following in both.
- Which algorithm has the highest accuracy? Is there model overfitting for this algorithm? Why or why not? Perform 10-fold cross-validation with this algorithm and report the accuracy, precision and recall.

the model with the highest accruacy was the Knearest neighbors model with K = (3, 5 and 9) all achieving perfect accuracy (no 10-fold validation yet), yes there was some over fitting but not too much, after performing the 10-fold CV the model achieved `CV Acc=0.973, Test Acc=0.967` these results are close to the origional output. see below for log table or all is avaible at the code block above.

Output logs for both vlaues of k after 10-fold CV:

k	test_acc	test_prec	test_rec	cv_acc	cv_prec	cv_rec	
6	13	0.966667	0.969697	0.966667	0.973333	0.977778	0.973333
7	15	0.966667	0.969697	0.966667	0.960000	0.966032	0.960000